

Quality gates for **AI-written** code.

What a gate can check on every pull request today, where it stops, and what we're building next.

A build log from the unslop team. Research numbers are third-party studies, sourced on each page.

Checks


Example



Anti-Slop Gatekeeper


1 critical slop finding

AI-attributed code carried more high-severity findings.



High-severity findings per file: 955 AI-attributed files vs. 955 matched human-control files. Same direction in JavaScript, Python and TypeScript.

“Improve this code”, ten times over, made it less secure.



Average vulnerabilities per sample when GPT-4o was asked to improve the same code again and again. Prompts that asked for security added new ones too, just fewer.

Today: a check on every pull request.

A critical finding fails the check. Warnings never block a merge. Every finding names its rule and, where it can, a one-click fix.

Checks

Example



Anti-Slop Gatekeeper

1 critical slop finding

unslop-gatekeeper

bot

Example



SEC-031 (CRITICAL)

An error is silently swallowed and the handler returns a success response anyway, so failures (including skipped authorization checks) look identical to successes.



Suggested fix (1-click commit)

Check name, titles and comment format are the product's own strings. Finding and code are example data.

What it doesn't do. Yet, or ever.

It is not fully deterministic.

119 rules from published research. 55 have a deterministic detector; the rest are checked by an LLM reviewer. Each finding says how it was verified.

The LLM review reads JS/TS only, for now.

Other languages get the deterministic pre-scanner.

It doesn't compare revisions yet.

Every push gets its own review. Failing a check because a revision got worse is on the roadmap.


It won't catch everything.

Nothing blocks 100% of AI vulnerabilities. We document the gaps.

Next: see the trend, not just the PR.

Planned: findings per repository over time, so you can see whether AI-assisted changes add debt faster than you remove it.

Concept sketch · example data · not built



Which metric belongs on this chart is an open question. Tell us what you would track.

Planned: one record per commit SHA. It lists which rules ran and what passed, so you can hand an auditor or an acquirer a receipt instead of a slide.

```
repository      acme/shop  
commit         3f9c2a17de...  
rules checked   119 (55 deterministic)  
critical        0  
warnings        3, reviewed by a human: 1  
not reviewed    0 files  
signature       ██████████
```

What it would never say: “this code is bug-free”.

If you run an engineering org: what would that record need to show before you trusted it?

We read every comment. The answers go into the spec.

Private beta, EU inference (Vertex AI, EU endpoint). Waitlist:
unslop.codes